

BIG DATA LANDSCAPE: NoSQL Systems

by Asst. Prof. Kulsawasd Jitkajornwanich, Ph.D.

`kulsawasd.ji@kmitl.ac.th`

[Adapted from slides and/or other materials from: Bengfort & Kim “Data Analytics with Hadoop”, White “Hadoop: The Definitive Guide”, Schutt & O’Neil “Doing Data Science”, Elmasri & Navathe “Fundamentals of Database Systems”, Cielen et al. “Introducing Data Science”, and Provost & Fawcett “Data Science for Business”]



Kulsawasd Jitkajornwanich, Ph.D.

Assistant Professor of Computer Science

- Received Master's and PhD degrees in Computer Science from The University of Texas at Arlington, TX, USA



Dissertation: “Analysis and Modeling Techniques for Geo-Spatial and Spatio-Temporal Datasets”



- Received Bachelor's degree (Hons.) in Computer Science from Chulalongkorn University, Thailand
- Conducting research/projects and academic services in the areas of database, NOSQL, big data analytics, social sciences and custom location-based applications



- Work Experiences:     

พระจอมเกล้าลาดกระบัง
King Mongkut's Institute of Technology Ladkrabang

Learning Objectives:

- 1) Understand concepts of nosql systems as well as other related terminology
- 2) To be able to communicate nosql concepts with the team effectively

Course Outline (part 2)


NoSQL System Fundamentals

- What is NoSQL? And NoSQL Background
- CAP Theorem
- Types of NoSQL systems
 - Some examples for each type of NoSQL

MongoDB: Document-Based NoSQL

- MongoDB Concepts & Data model
- Comparison to RDBMS
- CRUD operations
- Replication & Sharding
- Hands-on/Lab Exercises

Pyramid of Big Data Architecture Layers



Background

- Relational databases → **mainstay** of business
- Web-based applications caused spikes
 - explosion of social media sites (Facebook, Twitter) with large data needs
 - rise of cloud-based solutions such as Amazon S3 (simple storage solution, supporting ‘Data Lakes’)
- Hooking RDBMS to web-based application becomes trouble

Issues with *scaling up*

- Best way to provide ACID and rich query model is to have the dataset on a single machine
- Limits to *scaling up* (or *vertical scaling*: make a “single” machine more powerful) → dataset is just too big!
- *Scaling out* (or *horizontal scaling*: adding more smaller/cheaper servers) is a better choice
- Different approaches for horizontal scaling (multi-node database):
 - Distributed/parallel RDBMS
 - **NOSQL** (Non-RDBMS)

What is NOSQL?

- The Name:
 - Generally interpreted as “**Not Only SQL**”
 - The term NOSQL was introduced by *Carl Strozzi* in 1998 to name his file-based database
 - It was again re-introduced by *Eric Evans* when an event was organized to discuss open source distributed databases
 - Eric states that “... *but the whole point of seeking alternatives is that you need to solve a problem that relational databases are a bad fit for. ...*”



What is NOSQL?

- Key features (advantages):
 - non-relational
 - don't require schema
 - data are replicated to multiple nodes (so, identical & fault-tolerant) and can be partitioned:
 - down nodes easily replaced
 - no single point of failure
 - horizontal scalable
 - cheap, easy to implement (open-source)
 - massive write performance
 - fast key-value access



What is NOSQL?

- Disadvantages:
 - Don't fully support relational features
 - no join, group by, order by operations (EXCEPT within 'partitions' --table)
 - no referential integrity constraints across partitions
 - No declarative query language (e.g., SQL) → more programming
 - Relaxed ACID (see CAP theorem) → fewer guarantees
 - No easy integration with other applications that support SQL

What is NOSQL?

RDBMS	NoSQL
- User-friendly/mature features (join, group-by)	- New tech./technical background required
- Strong consistency (<i>ACID</i>)	- Relaxed consistency (<i>BASE</i>)
- Can be slow (if > PB)	- High availability
- Structured (relational) data	- Dynamic/complex data structure
- Vertical scaling (“scaling up”)	- Horizontal scaling (“scaling out”)

Who is using them?




3 major papers for NOSQL


- Three major papers were the “seeds” of the NOSQL movement:
 - BigTable (Google)
 - DynamoDB (Amazon)
 - Ring partition and replication
 - Gossip protocol (discovery and error detection)
 - Distributed key-value data stores
 - Eventual consistency
- **CAP Theorem**

CAP Theorem

- Suppose three properties of a distributed system (w/ data replication)
 - **C**onsistency:
 - always return the most up-to-date values
 - **A**vailability:
 - always respond (either failed or succeeded)
 - **P**artition tolerance:
 - can still operate even if there is a network fault resulting in two or more partitions (allowing some messages to get lost)




CAP Theorem



All client always have the
same view of the data


CAP Theorem



Each client **always** can read and write.

CAP Theorem

A system can **continue to operate** in the presence of a network partitions



CAP Theorem

- A **consistency** model determines rules for visibility and apparent order of updates
- Example:
 - Row X is replicated on nodes M and N
 - Client A writes row X to node N
 - Some period of time t elapses
 - Client B reads row X from node M
 - **Does client B see the write from client A?**
 - Consistency is a continuum with tradeoffs
 - **For NOSQL, the answer would be: “maybe”**
 - CAP theorem states: *“strong consistency can't be achieved at the same time as availability and partition-tolerance”*

NOSQL categories

1. Key-value

- Example: DynamoDB, Voldermort, Scalaris

2. Document-based

- Example: MongoDB, CouchDB

3. Column-based

- Example: BigTable, Cassandra, Hbase

4. Graph-based

- Example: Neo4J, InfoGrid
- “No-schema” is a common characteristic of most NOSQL storage systems
- Provide “flexible” data types (eg, JSON, XML, ...)
- Rigid constraints on index creation
 - Mostly ONE index and SINGLE attribute (not composite) are allowed

NOSQL categories

1. Key-value

- Example: DynamoDB, Voldermort, Scalaris

2. Document-based

- Example: MongoDB, CouchDB

3. Column-based


- Example: BigTable, Cassandra, Hbase

4. Graph-based

- Example: Neo4J, InfoGrid
- “No-schema” is a common characteristic of most NOSQL storage systems
- Provide “flexible” data types (eg, JSON, XML, ...)
- Rigid constraints on index creation
 - Mostly ONE index and SINGLE attribute (not composite) are allowed

Document-based

- Can model more complex objects
- Inspired by Lotus Notes
- Data model: collection of documents
- **Indexing:** automatically indexed on ONE selected unique attribute (hash OR range (partition))
- Document: JSON (**J**ava**S**cript **O**bject **N**otation) is a data model supporting objects, records, structs, lists, array, maps, dates, Boolean (with **nesting**), XML, other more complex semi-structures.
- Comparison to key-value NOSQL:
 - Slower but more flexible since: 1) condition on an un-indexed attribute can be specified in the query, 2) normalization concept is allowed, 3) only ONE attribute can be indexed, and 4) more complex objects can be modelled.



Document-based

- Example: (MongoDB) document

- {Name:"Jaroslav",

Address:"Malostranske nám. 25, 118 00 Praha 1",

Grandchildren: {Claire: "7", Barbara: "6", "Magda: "3", "Kirsten: "1", "Otis: "3", Richard: "1"}

Phones: ["123-456-7890", "234-567-8963"]

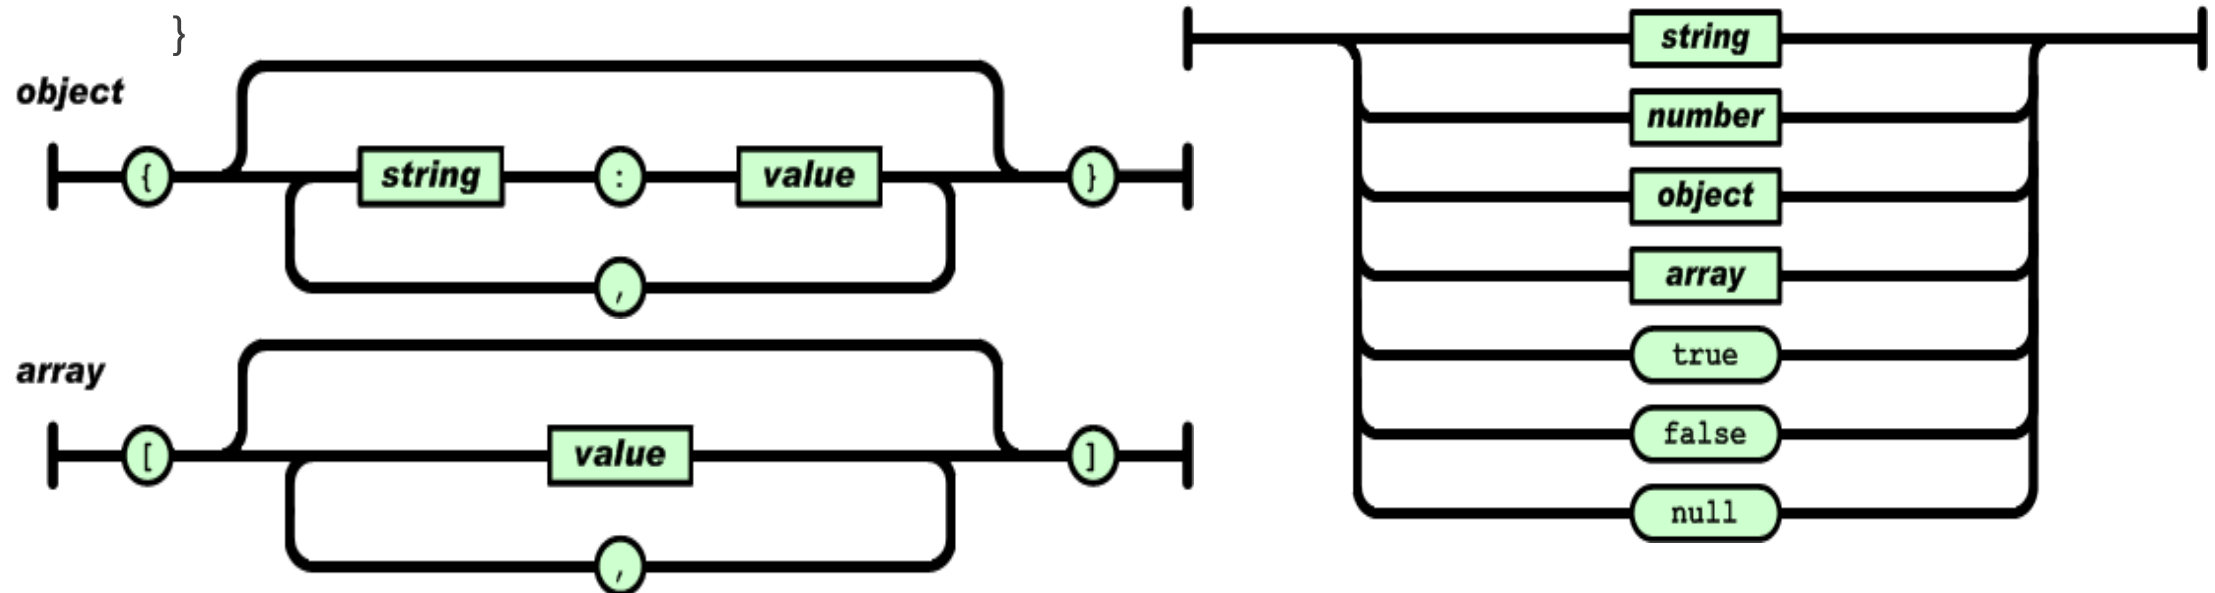


Figure 3.6

One possible database state for the COMPANY relational database schema.

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

Continued next page...

Document-based

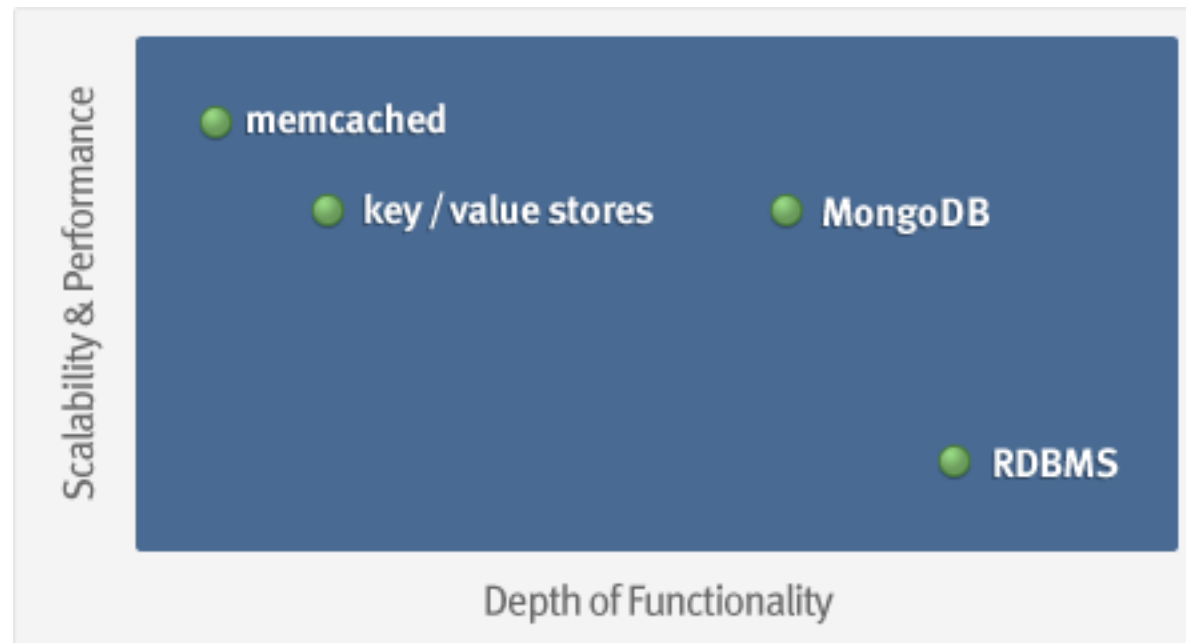
Name	Producer	Data model	Querying
MongoDB	10gen	object-structured documents stored in collections; each object has a primary key called ObjectId	manipulations with objects in collections (find object or objects via simple selections and logical expressions, delete, update,)
Couchbase	Couchbase ¹	document as a list of named (structured) items (JSON document)	by key and key range, views via Javascript and MapReduce

MongoDB: Document-based NoSQL

- stands for humongous
- created by 10gen
- open source (Implemented in C++)
- document-oriented database designed with
 - 1) scalability
 - 2) developer agility in mind
- Instead of storing data in tables and rows as we would do with a relational database, in MongoDB we store **BSON** (binary representation of **JSON**) documents with dynamic schemas (schema-less).

MongoDB: Background

- Goal:
 - bridge the gap between key-value stores (which are fast and scalable) and relational databases (which have rich functionality).



MongoDB: Background

- Key features:
 - scale “horizontally” over commodity hardware
 - support RDBMS features:
 - Ad hoc queries
 - Fully featured indexes
 - Single-field
 - Compound
 - Multi-key
 - Geospatial
 - Text
 - Aggregation operations (through *map-reduce*)

Data model

- DB ← Collections ← Documents
- **Collection** = set of “related” documents sharing “common” indexes
- **Primary key** (`object_id`) is automatically created and indexed for a document
 - Help narrow search span; otherwise MongoDB needs to scan all documents in every shards

Data model

- Schema design:
 - Embedding
 - nesting of objects and arrays inside a BSON document (pre-joined) → fast
 - Link (referencing)
 - references between documents → requires follow-up query
- Principle:
 - “embedding when you can, link when you must”

Comparison to RDBMS

RDBMS	MongoDB
Database	Database
Table, View	Collection
Row	Document (BSON → JSON)
Column	Field
Index	Index
Join	Embedding/Referencing
Foreign Key	Referencing
Primary Key	ObjectID

CRUD operations

- **CRUD**

- stands for **C**reate, **R**ead, **U**ppdate, and **D**eleate
- used for reading and manipulating data

- Example: JSON documents and collection

```

{
  name: "sue",
  age: 26,
  status: "A",
  groups: [ "news", "sports" ]
}

```

← field: value
 ← field: value
 ← field: value
 ← field: value

```

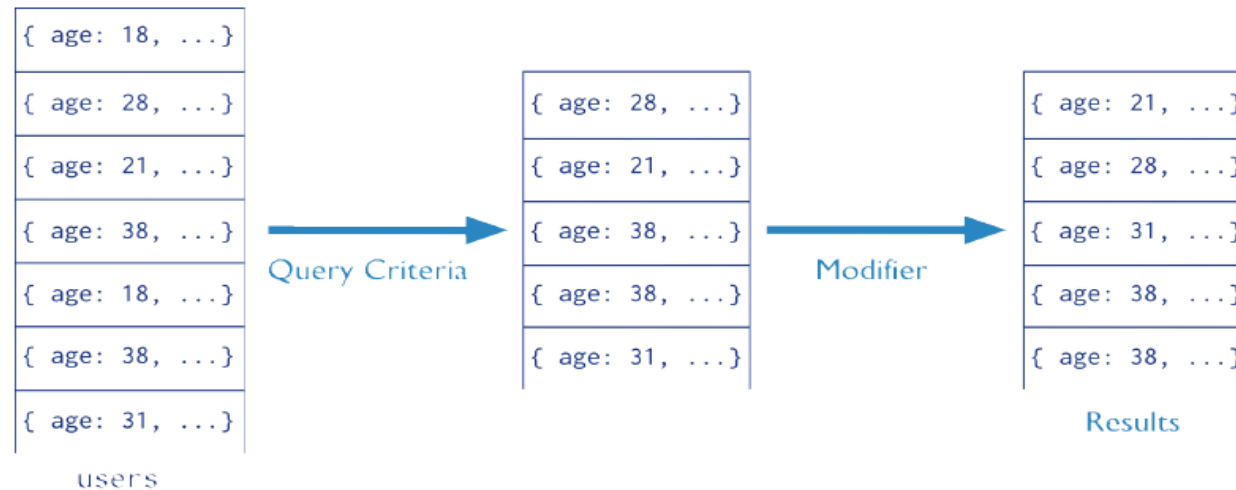
{
  name: "al",
  age: 18,
  status: "D",
  groups: [ "politics", "news" ]
}

```

CRUD operations

- target *a specific collection at a time*
 - NOT** support JOIN operations


Collection Query Criteria Modifier
 db.users.find({ age: { \$gt: 18 } }).sort({age: 1 })



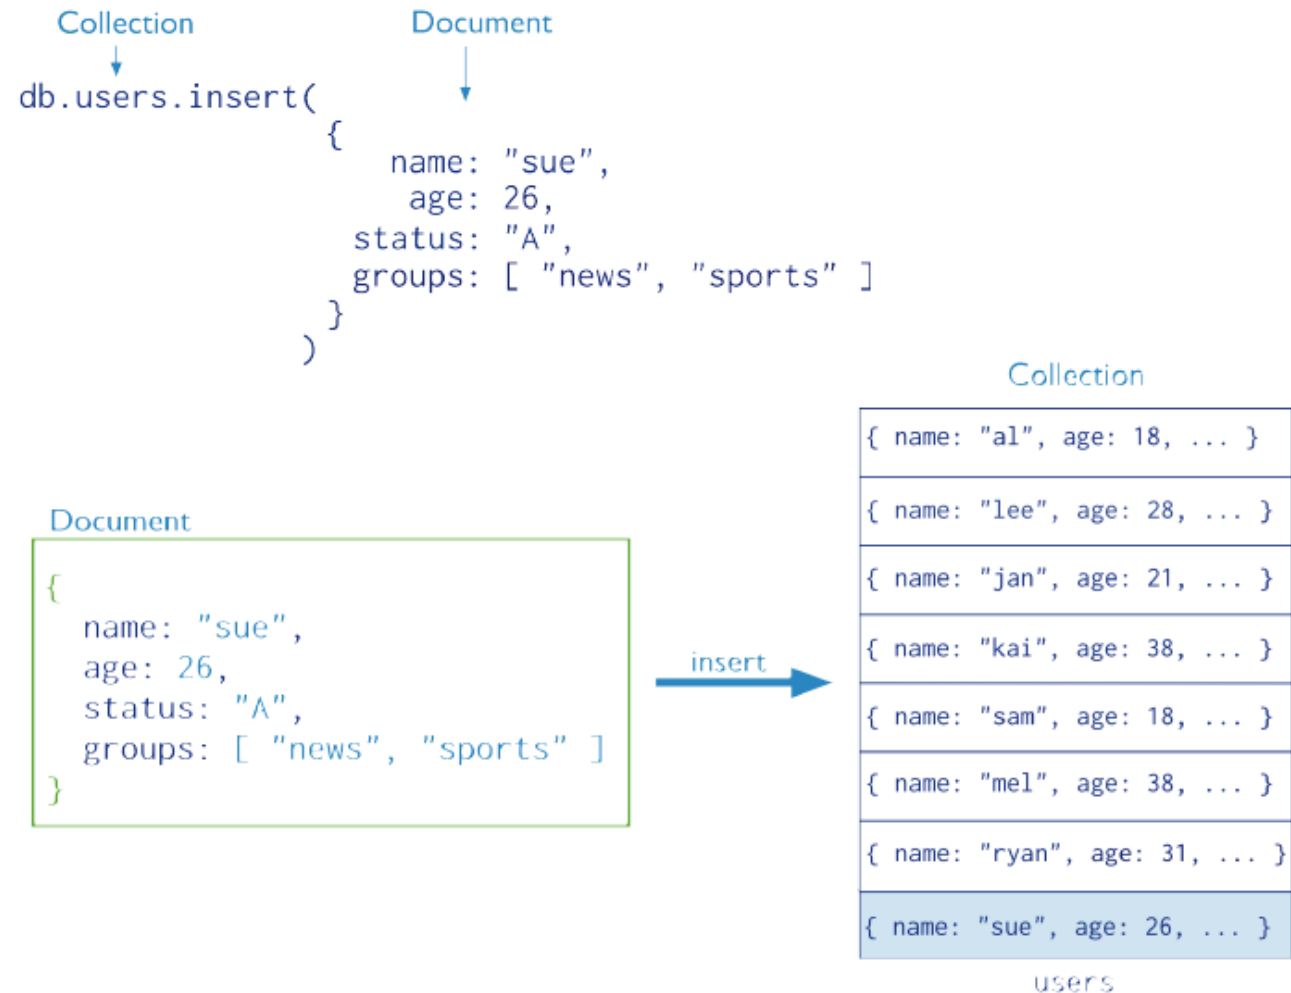
CRUD operations

- More example:

Collection Query Criteria Projection
`db.users.find({ age: 18 }, { name: 1, _id: 0 })`



CRUD operations



CRUD operations

```
SELECT _id, name, address  
FROM   users  
WHERE  age > 18  
LIMIT 5
```

← projection
← table
← select criteria
← cursor modifier



```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

► Note: Results always include object id (*id*) unless explicitly specify (id:0)

- Results always include object id (*id*) unless explicitly specify (id:0)

CRUD operations

- Create

- db.collection.insert(<document>)
- db.collection.update(<query cond>, <update>, { upsert: true })

- Read

- db.collection.find(<query cond>, <projection>)

- Update

- db.collection.update(<query cond>, <update>, <options>)

- Delete

- db.collection.remove(<query cond>, <justOne>)

Wait! Why ‘Just One’? What about SQL system?

Conclusion

- NOSQL databases cover only a part of data-intensive cloud applications (mainly Web applications)
- Problems with cloud computing:
 - SaaS (**S**oftware **as a S**ervice or on-demand software) applications require enterprise-level functionality, including ACID transactions, security, and other features associated with commercial RDBMS technology, i.e. NOSQL should not be the only option in the cloud
 - Hybrid solutions:
 - Voldemort with MySQL as one of storage backend
 - deal with NOSQL data as semi-structured data
 - integrating RDBMS and NOSQL via SQL/XML

Conclusion

- next generation of highly scalable and elastic RDBMS: *NewSQL databases*
 - they are designed to scale out horizontally on shared nothing machines,
 - still provide ACID guarantees,
 - applications interact with the database primarily using SQL,
 - the system employs a lock-free concurrency control scheme to avoid user shut down,
 - the system provides higher performance than available from the traditional systems.
- Examples: MySQL Cluster (most mature solution), VoltDB, Clustrix, ScalArc, etc.



Any Questions?